



How TradePilot Works

A robot team that watches 889 companies at once, decides which two are worth buying today, and never blinks. Written so that it makes sense even if you have never bought a share in your life.

PREPARED BY

Soumya Swain

DATE

25 August 2026

STATUS

Paper trading — real money Wednesday

First, the five words you need

Everything in this document is built out of five ideas. If you understand these, you will understand the whole system.

SHARE — A company splits itself into millions of tiny pieces. One piece is a share. If you own one share of a company worth ₹400 a share, you own ₹400 of that company. Buy it at ₹400, sell it at ₹406, and you made ₹6.

THE MARKET — A giant public auction where about 900 Indian companies have their shares bought and sold. It runs 9:15am to 3:30pm, Monday to Friday. Prices change several times *per second* because thousands of people are bidding at once.

THE TOLL — Every time you buy and then sell, you pay fees and government taxes. For us that is 0.106% of the money involved. Trade ₹10,000 and it costs you about ₹10.60 — *whether you win or lose*. This one number is the villain of the entire story.

A LEVEL — A specific price that lots of people are watching, like yesterday's lowest price. Prices tend to do something interesting when they arrive at one — bounce off it, or break through it. Levels are where our system pays attention.

PAPER TRADING — Playing the whole game with fake money to see whether the plan works, before risking real money. Everything in this document is currently paper. Real money starts Wednesday, with ₹3,000.

Why the toll changes everything


Most explanations of trading talk about being right more often than you are wrong. That is not the hard part.

The hard part is that **you must be right by more than 0.106% every single time** — otherwise you go slowly broke while technically "winning" more often than you lose. Being right is not enough. You have to be right by *enough*.

Ten famous ideas. All ten lost money.

We tested ten well-known trading ideas against real market history — **145,500 simulated trades** across 201 companies. Not opinions: measurements.

The best idea earned 0.051% per trade. The toll is 0.106%. So the single best-performing famous strategy still lost about half the toll on every trade it took. That is the problem TradePilot exists to solve, and it is why the system spends most of its effort *refusing* trades.



What each strategy earned on an average trade. The red line is the toll — every bar has to reach it to break even, and not one of them does. The best famous idea falls short by 0.015% per trade; trend following falls short by 0.145%. That picture is the entire difficulty of this business.

So what is left?

If the famous ideas do not work, the answer is not a cleverer idea. It is to be far more selective about *when* to trade at all — and to find something that holds up when everything else is tested honestly. We found exactly one such thing, and it is on page 8.

What the system actually does

Imagine a huge marketplace with 889 shops. Every shop's price changes every few seconds. You have enough money to buy from two shops today. Which two, and when exactly?

No human can watch 889 price boards at once. So TradePilot uses two different kinds of worker, because the job is really two jobs:

4 Scouts

Broad and shallow. They walk the entire market once every minute and glance at all 889 shops. They cannot study any one shop closely — but nothing escapes them.

every 60 seconds
889 companies
4 numbers each

20 Watchers

Narrow and deep. Each one stands in front of exactly one shop and stares at its price board without blinking. It knows that shop perfectly — and is blind to the other 888.

every price change
1 company each
thousands of looks/min

1 Judge


Slow and careful. Only called when a watcher raises its hand. Looks at the actual chart, asks six questions, and either approves a trade or refuses it.

only on request
~seconds per decision
can say no

The scouts decide **what to look at**. The watchers decide **when something happened**. The judge decides **whether to actually buy**. None of them can do the others' job, and that is the point — a scout that tried to watch closely would see almost nothing, and a watcher that tried to scan everything would notice nothing in time.

889 companies in, 2 trades out

Every layer below is a filter, and each one throws work away. These numbers come from a real scan on 24 August 2026 — they are not an illustration.



Read it top to bottom. The whole market enters at the top; two trades leave at the bottom. The system's main job is saying no.

STAGE	WHAT GETS REMOVED	WHY
938 → 889	Funds and baskets	These are bundles of many companies, not companies. Five gold funds once filled our shortlist and looked like five separate opinions — they were one opinion counted five times.
889 → 437	Too expensive, too quiet, or frozen	Shares costing more than ₹600 we cannot size properly. Companies barely traded today would move against us when we buy. And a share frozen at its daily limit has nobody to sell it to us.
437 → 60	Nothing interesting today	At least one of the four scouts has to find a reason.
60 → 20	Not the strongest reasons	We only have twenty watchers.
20 → 2	The moment never arrived	Most watched companies never trigger anything worth acting on. That is normal and expected.

The four scouts, and what each one hunts

Four scouts, not one — and deliberately four *different* ones. If all four looked for the same thing, we would have one scout at four times the cost. Each one is built to notice what the other three physically cannot see.

1 - Trend

Looks for companies that have been climbing steadily for weeks — but have *paused* today. Like a strong runner catching their breath.

rising 20 & 50 day averages
higher highs, higher lows
NOT up more than 4% today

2 - Flow

Looks for unusually heavy buying and selling versus that company's own normal. Doesn't care which direction. This is the "something is happening here" scout.

at least 1.8× normal volume
measured as a rate, not a total
fires before price moves

3 - Level

Looks for prices sitting right on top of an important number — a 52-week high, yesterday's close, a round number like ₹500.

within 0.6% of a level
52-week high scores highest
middle of nowhere = ignored

4 - Reversal

Looks for good companies that have been beaten down for five days and are just starting to bounce. The Trend scout rejects every single one of these.

down more than 3% over 5 days
bounce already starting
still above its 200-day average

The Reversal scout is the clearest illustration of why four are needed. It hunts companies that have been *falling*. The Trend scout rejects every one of those automatically. Neither is wrong — they are looking for different opportunities, and a single combined scout would simply miss one of them.

How a scout turns a company into a score

Each scout gives every company a score between 0 and 1. It starts from a base number and then adds or subtracts points for specific, measurable facts — no opinions anywhere. Here is the Trend scout in full. The other three work the same way.

WHAT IT CHECKS	EFFECT ON SCORE	WHY
Price above its 20-day and 50-day average	required	If this fails, it isn't a trend at all
Today's move outside -2% to +4%	rejected outright	Crashing or already exploded — both are bad entries
Starting score	0.30	—
Making higher highs AND higher lows	+0.25	The textbook shape of a healthy climb
Above its 200-day average	+0.10	Healthy over the long run, not just recently
Up more than 5% over 20 days	+0.10	The climb is real, not noise
Sitting right on its 20-day average	+0.20	This is the pause we want to buy
Already up more than 3% today	-0.30	We are late. Chasing loses money.

Why we punish a stock for going up

This surprises people. Surely a stock rising fast is a *good* sign?


Not within a single day. We measured this over 145,500 trades: buying in the direction of the day's move **lost** 0.039%, while buying against it gained 0.003%. Within one day, Indian shares tend to snap back rather than run. So we look for strong companies that are *resting*, never ones already flying.

This is a good example of how the system is built: a rule that sounds wrong, kept because the measurement says so.

The one thing that actually worked

We tested ten well-known trading ideas. Nine of them were somewhere between useless and harmful. But one finding survived every test we threw at it, and it is the reason the whole system is built the way it is.

When separate scouts agree, the odds improve — every time



Each point is thousands of real trades. On the left: only one scout was interested, and those trades lost badly. On the right: seven independent signals agreed, and those trades made money. The line climbs without a single dip — that consistency is what makes it trustworthy.

Think about why this makes sense. If one friend tells you a shop is worth visiting, they might be wrong, or biased, or just excited. If four friends who went there for *completely different reasons* all say it independently — a busy queue, a discount sign, a good smell, a familiar face — you should probably go.

That is exactly what the scout board does. The four scouts never talk to each other. They look at different things entirely. When two or three of them independently land on the same company, that agreement is worth more than any single scout being extremely confident.

Agreement, made into a rule

Because agreement is the only thing we proved, the system is built to reward it directly — and getting that right took a correction.

So agreement beats enthusiasm — and we had to say so explicitly

We originally combined the four scores into one number with a formula, hoping agreement would come out naturally. It nearly worked. Then we checked, and found a case where a company all *three* scouts liked was rejected in favour of one that two scouts liked — losing by 0.009 points.

A formula was quietly throwing away the one thing we had actually proved. So we stopped being clever and wrote the rule down plainly: **more scouts agreeing wins, full stop.**

SCOUTS AGREEING	SCORE RANGE	COMPANIES ON 24 AUG	WHAT IT MEANS
3	0.931 – 0.964	3	Rare and strong. Watch these.
2	0.553 – 0.850	26	Worth a watcher.
1	0.492 – 0.600	31	Interesting, not convincing.

Notice the gaps between those bands. A company cannot drift from "two scouts" to "three scouts" by a rounding error — a third scout has to genuinely find a reason. That gap is deliberate: it means the system only changes its mind when something real has changed, not when a number wobbles.


The agent floor, in detail

This is the newest part of the system and the part people ask about most. Twenty independent agents, each assigned one company, each watching it continuously for the entire trading day.

Why not just have one clever AI watch everything?

Because of speed. An AI that "thinks" takes a few seconds to answer. A price move worth catching also takes a few seconds. By the time the AI finished thinking, the moment would be gone — every time.

So we split the job the way a real trader actually works. A trader does not re-analyse everything every second. They **decide their important prices in the morning**, then simply **watch for those prices to arrive**. That is two very different activities, and only the second one needs to be fast.



One agent. On the left, the small set of facts it memorises at 9:16am. On the right, the only five events that make it interrupt anyone. Everything else it sees, it silently ignores.

Fast eyes, slow judgement

The whole design rests on separating two things people usually lump together: *noticing* that something happened, and *deciding* what to do about it.

LAYER	SPEED	WHAT IT DOES	RUNS HOW OFTEN
Fast — plain code	microseconds	Compares the live price against the agent's memorised levels. No thinking, just arithmetic.	every price change
Slow — judgement	seconds	Reads the actual chart, asks six questions, approves or refuses the trade.	only when a hand goes up

The point of the split

Thinking is expensive and slow. Arithmetic is free and instant. By letting cheap arithmetic decide *when* something matters, and spending expensive judgement only on those rare moments, we can afford twenty agents watching continuously — and still respond faster than a human could.

The five things that make an agent speak

TRIGGER	IN PLAIN WORDS	EXACT THRESHOLD
Level touch	Price has arrived at one of the important prices it memorised	within 0.08%
Sweep and reclaim	Price dipped below an important price and jumped straight back above it	dip >0.05%, then recover
Fast move	Price is moving unusually quickly right now	0.25% in a minute
Volume burst	A sudden rush of buying and selling	3× the recent rate
Invalidation	A trade we already hold has gone wrong	instant, highest priority

After any one of these fires, that agent stays quiet about the same kind of event for three minutes. Without that pause, a single jumpy company could raise its hand hundreds of times and drown out the other nineteen.

One agent's working day

- 1 9:00am — Homework.** Before the market opens, the system studies about 320 companies: their averages, their typical daily swing, their 52-week high and low.
- 2 9:16am — Assignment.** The scouts publish their ranked list. The top 20 companies each get an agent. Each agent memorises four or five important prices for its company — yesterday's high and low, today's average, the nearest round number.
- 3 All day — Watching.** The agent receives every single price change from the exchange and compares it to its memorised levels. Thousands of comparisons per minute. It does not get bored, distracted or tired.
- 4 Whenever it matters — Raising a hand.** If one of the five trigger conditions is met, the agent escalates: it records what happened, at what price, and at what time, and asks for judgement.
- 5 Every 2 minutes — Possible reassignment.** The scouts are still scanning. If a much better company appears, an *idle* agent is moved to it.
- 6 3:30pm — Reporting.** Every escalation, every reassignment, and every near-miss is written to a file so we can check the next morning whether the agents were actually useful.

The three rules that stop agents wandering off

Moving agents around sounds obviously good, but done carelessly it destroys the system. Three rules prevent that:

Never abandon a trade

An agent holding an open position cannot be moved, no matter how attractive another company looks. You do not walk away from money you have already committed.

Ten minutes minimum

An agent must watch a company for at least 10 minutes before it can be moved. Two of its five triggers need 30 and 60 price-changes of history — an agent moved constantly would be permanently blind.

Four moves maximum

At most four agents change company in any one cycle, and the newcomer must be clearly better, not marginally. Otherwise the floor would churn all day and settle on nothing.

What we are honestly unsure about

Nobody has run this live yet. Today is the first real session. We do not know whether 20 agents will produce 15 useful alerts or 500 useless ones — so the system records every alert *and* every alert it nearly raised. If the settings are wrong, the log will say so, and we will change them rather than guess again.

Watch one trade happen, step by step

This is the clearest way to understand the whole system: follow a single company through one morning, from the agent's first glance to the order.



A company's price through one morning. The two dashed lines are levels the agent memorised at 9:16am. Everything the agent does is a reaction to those lines.

Why "sweep and reclaim" is worth waking someone up for

Many people leave a standing instruction with their broker: "sell my shares automatically if the price drops below ₹411." When the price does dip under ₹411, all of those instructions fire at once and shove the price down briefly.

If the price then climbs straight back above ₹411, it tells you something specific: those forced sellers are finished, and buyers were waiting underneath the whole time. The dip on its own means nothing. **The recovery is the signal** — which is exactly why the agent stays silent at 11:02 and only speaks at 11:04.

The same trade, minute by minute

Here is the identical sequence written out, so you can see how little the agent actually does — and how much of the day it spends deliberately silent.

TIME	WHAT HAPPENS	WHAT THE AGENT DOES
9:16	Assigned to this company by the scouts — Trend and Level both flagged it	Memorises: yesterday's low ₹411.00, today's average ₹414.50
9:30–10:40	Price drifts up and down between ₹411 and ₹416	Nothing. Watches ~40,000 price changes, stays silent.
11:02	Price falls <i>below</i> ₹411 — under yesterday's low	Notes it. Still silent — a break alone is not the signal.
11:04	Price jumps straight back <i>above</i> ₹411	Raises its hand. This is the "sweep and reclaim" pattern.
11:04	Judgement is called	Six questions asked. Where is price? At a level. What is the structure? Uptrend. What broke? Nothing. Where am I wrong? Below ₹410.90.
11:05	Approved — a trade card is sent to a phone	Buy 14 shares at ₹414.40. Stop at ₹411.10. Target ₹420.60.
14:45	Either the target, the stop, or the clock ends it	Position closed. Nothing is ever held overnight.

Look at what the agent did NOT do

Between 9:30 and 10:40 it examined roughly forty thousand price changes and said nothing at all. It did not get bored, did not talk itself into a trade, and did not lower its standards because the morning was quiet.

That patience is the actual product. Anyone can find reasons to buy. The difficult part — the part the toll punishes — is sitting still for four hours and acting only when a specific, pre-defined thing happens.

The six questions before any money moves

An agent raising its hand is not permission to trade. Every candidate must survive six questions in order. Fail the first two and the answer is simply no.

- 1 **Where is the price right now?** Name the nearest important price above and below, with the distance to each. If we are in the middle of nowhere, there is no trade — the edge only exists at levels.
- 2 **What shape is the chart in?** Climbing, falling, or going sideways — and name the exact price that would prove it has changed. If you cannot name that price, there is no trade.
- 3 **What specific thing just happened?** "Dipped below yesterday's low and recovered." Not "it looks good." A feeling is not a trigger.
- 4 **What would prove me wrong?** An exact price, decided *before* buying. That price becomes the automatic exit, and the gap between it and the entry is exactly how much we can lose.
- 5 **How many separate reasons agree?** Count them. Fewer than four and the profit will not cover the toll.
- 6 **How much is there to win?** Distance to the next obstacle, divided by the amount at risk. Below 1.5 times, skip it.

A confidence figure between 0 and 1 is stated, and anything under 0.6 is skipped. **A day with no trades at all is a perfectly good outcome** — and often the correct one.

Question 4 is the one that protects the money

Deciding the exit price *before* buying sounds like a technicality. It is the single most important rule in the document. It converts "I hope this goes up" into "I lose exactly ₹X if I am wrong" — a known, bounded, survivable number.

Everyone who blows up an account skipped question 4.

Where the money actually goes

We start with ₹3,000. For trades opened and closed on the same day the broker lends four times that, so we control about ₹12,000 — split into two positions of ₹6,000 each.

POSITIONS	EACH	MOST EXPENSIVE SHARE WE CAN BUY	COMPANIES AVAILABLE	COST PER TRADE
1	₹12,000	₹1,200	100%	0.1060%
2 ← chosen	₹6,000	₹600	82%	0.1060%
3	₹4,000	₹400	58%	0.1060%

Two positions cost exactly the same as one

Look at the last column: splitting the money changes nothing about the cost, because every fee is a straight percentage. So two positions are **free diversification** — and they double how much we learn per day, which matters when the plan is to judge this over ten sessions.

Three would be free too, but the share price limit drops to ₹400 and we would lose 42% of the companies we can choose from. Two is the sweet spot.

Why the price limit exists at all

With ₹6,000 you can buy 10 shares of a ₹600 company — or 2 shares of a ₹2,500 company. With only 2 shares, you cannot size a position properly; you are stuck jumping in huge steps. We insist on at least 10 shares, and that single rule is what sets the ₹600 ceiling.

The second experiment: options

Alongside buying shares, we are testing something different — and riskier — with a small, fixed amount.

AN OPTION — A ticket that says "I may buy the market at ₹24,250 before next Tuesday". If the market climbs well past that, the ticket becomes valuable. If it doesn't, the ticket expires and is worth exactly nothing. You can lose 100% of what you paid — but you can never lose more than that.

Options are attractive because a small amount of money can produce a large gain. They are dangerous because **the most likely single outcome is losing everything you put in**. That is why this lane is capped, is being tested on paper first, and must pass a gate before it ever sees real money.

BUDGET	CLOSEST TICKET WE CAN AFFORD	COST	FEES AS % OF THE TRADE
₹1,000	far away — needs a huge move	₹725	6.5%
₹3,000	very close to the current price	₹2,532	1.9%

This is the clearest argument for raising the amount from ₹1,000 to ₹3,000. At ₹1,000 we could only afford tickets so far from the current price that they were effectively lottery tickets. At ₹3,000 we can buy one that responds to ordinary market moves — and the fees drop from 6.5% to 1.9% of the trade.

Tomorrow is expiry day — so extra rules apply

Options have a deadline. Tomorrow is the deadline for this week's batch. On deadline day a ticket that is not clearly winning loses value very fast, and by 3:30pm it is worthless with certainty.

Our budget forces us onto this day — it is the only one where ₹3,000 buys a ticket close to the current price. So we guard it: **no new tickets bought after 1:00pm**, everything sold by **3:00pm**, and we only buy tickets that plenty of other people are trading, so there is always someone to sell back to.

Why this lane is capped, and gated

The options experiment cannot spend real money until it has produced **eight paper trades that are profitable after fees**. It has produced zero so far. That gate is not a formality — six earlier ideas in this project were killed by exactly this kind of pre-agreed test, and killing them was the right outcome each time.

What we know, and what we don't

The most useful thing in this document is the honest list of what remains unproven. Every row below has a specific test attached, decided in advance, so that we cannot later argue ourselves into a favourable interpretation.

QUESTION	HOW IT GETS ANSWERED	WHERE WE ARE
Do 20 agents produce useful alerts?	Count and quality of alerts in today's log	first live run
Should agents move companies mid-day?	Did companies moved to beat the ones left behind?	measured, not yet known
Does reading charts beat pure arithmetic?	20 trades, profitable after costs, beating random picks	open
Do options work at this size?	8 paper trades, profitable after ₹47 of fees	0 of 8
Do real orders fill at the price we expect?	10 real sessions, fills within 0.15% of the plan	starts Wednesday

The honest summary

We have built a system that watches the whole market properly, filters it down with rules we can defend, and refuses far more trades than it takes. We have proved that most popular trading ideas do not survive costs, and we have found one that does: independent agreement.

We have not yet proved that the finished system makes money. That is what today, and the ten sessions after it, are for. Every number in this document came from a measurement, and where we do not know something, this document says so.

If someone shows you a trading system with no page like this one, that is the page they chose not to write.