



Session Report

Wednesday 26 August 2026 — the day the agents started trading on their own, and the day their first result was falsified within two charts.

PREPARED BY

Soumya Swain

SESSION

2026-08-26

CAPITAL AT RISK

None — paper throughout

The day at a glance

Six commits. Real money deliberately withheld. The escalation-to-trade loop closed for the first time — then falsified within two charts, and undermined again by a bug found at 14:15.

31

PAPER TRADES

-₹501

NET

32%

WIN RATE

16

DECLINED

0

DATA GAPS

Real money was held back, and every gate agreed

Wednesday was the day ₹3,000 of real cash was meant to go in. It did not, and not out of nerves — all four pre-agreed gates said no.

BLOCKER	STATE	WHY IT MATTERS
Account funding	₹1,000 of ₹3,000	The top-up had not landed. The plan is undefined at the size it was designed for.
Options gate	0 of 8	Eight closed paper cards, net positive after fees, was the agreed bar.
Equity lane	0 cards ever	It has never produced a trade card on any day.
Infrastructure	53% down Tuesday	The floor died at 12:49 and never recovered.

And the measurement that reframed the question

Tuesday's 1,563 alerts were scored against a control of random minutes in the same stock. Every trigger showed **volatility** afterwards and **no direction** — signed returns -0.014%, 45.6% of moves up, all inside the 0.106% cost of trading. What we have built is an excellent *attention allocator*. It says where something is about to happen, not which way.

09:13 – three changes, all from Tuesday's numbers

- 1 Volume-burst alert retired**
5b11ed1
Fired **540 times** — more than any other — with a lift of **-0.014pp** against the control. Loudest, and worse than random. Disabled rather than deleted so the call stays reversible.
- 2 Sweep-and-reclaim loosened**
Our strongest signal by **3x** (+0.278pp) and our *quietest* at 118 fires — the wrong way round, so we were missing valid instances. Pierce threshold 5.0 → 3.0 bps, lookback 30 → 45 ticks.
- 3 The restart path – the bug that cost 2h40m**
Tuesday's watchdog detected the dead floor correctly and restarted it twice. Both relaunches were dead within ten minutes having written *nothing*, not even an error.
The cause: the watchdog is itself a scheduled job, and macOS reaps the job's whole process group when the script exits — so a background child cannot outlive its parent. The watchdog logged "relaunched"; the floor was already gone. It now restarts the floor's own scheduled job, and **verifies twelve seconds later that a process actually exists**.

By 10:41 the changes had done what the data predicted

Alert rate 8.8/min → **6.5/min**. Sweep-and-reclaim up **71%**. **Zero data gaps**, where Tuesday had four by the same hour and was already 36 minutes blind. Five reassignments against zero in Tuesday's entire first run.

11:55 – positions read zero because none could exist

The console showed 0 open positions. Not a display bug: nothing in the floor had ever been able to take one.

`self.position` was assigned exactly once — to nothing — and never again. Three consequences, none cosmetic:

- **Two of five triggers were unreachable code.** The stop-loss and take-profit checks sat behind that guard and had never once been able to run.
- **A safety rule I had called load-bearing protected nothing.** "An agent holding a position is never reassigned" — there had never been a position to abandon.
- **The console faithfully reported a number that could only be zero**, which reads as "no trades today" rather than "this path does not exist".

12:09 – the loop closed, with no approval gate

RULE	SETTING	WHY
Trigger	sweep & reclaim only	The only alert that implies a <i>direction</i>
Agreement	2+ scouts	Confluence is our one finding that survived falsification
Stop	the swept low	Structural — if price returns below what it reclaimed, the story is false
Target	1.5R	Below this the toll eats the trade
Size / window	₹6,000 • 09:30–14:30	Max 5 concurrent, squared off 15:15
Fees	0.106% round trip	Our measured cost, so net is net

Stated plainly

This is **not believed to be profitable**. It exists so that acting on an alert produces *evidence* instead of a guess. Every refused trade is logged with its reason, so a quiet day can never hide a threshold silently rejecting everything.

12:13 – looking at the charts falsified the rule

I had the system draw the actual charts behind two of its own live signals and read them the way a person would. Both were no-trades, for the same structural reason.

STOCK	WHAT THE CHART SHOWED	WHY THE RULE WAS WRONG
UDS	Coiled in a narrow range all session, crossing the average price repeatedly	A "dip and recovery" through a line price has crossed twenty times is not a trap being sprung
QUESS	A textbook decline — each peak lower than the last: 374, 372, 370, 368	The rule would have bought into a steady fall , because a bounce back to the average satisfies the arithmetic

An average can never be a place where orders rest

Half of any day is below that day's average price — that is what "average" means. Price crosses it constantly, in every stock. So "dipped below the average and came back" is like saying "*the temperature fell below today's average and then rose*". Of course it did.

Nobody parks a standing order at an average. They park them at yesterday's low, today's low, at round numbers. **Those can be swept. An average cannot.**

Quantified before changing anything: signals on the average were **33% of Tuesday's and 42% of Wednesday's**. Entries are now restricted to real levels where orders genuinely rest — cutting the largest single slice of candidates, the slice the pictures showed was worthless.

This is exactly what the chart-reading lane was for

The rule can only check arithmetic. Looking at the picture catches the case where the arithmetic is technically true and completely meaningless. It took two charts.

14:08 – the first results, and they are losing

METRIC	VALUE	READING
Closed trades	31	First ever from the floor
Win rate	32% (10/31)	Losing more often than winning
Average win	₹+27.45	And losing <i>bigger</i> than winning — the worse of the two
Average loss	₹-36.95	
Gross	₹-309.45	Losing before fees are applied at all
Fees	₹191.87	31 round trips at 0.106%
Net	₹-501.32	
Median hold	3.2 min	Very short — the stop sits inside the noise

The shape of it: **21 stops against 10 targets**. The stop at the swept low is structurally correct but tight, so ordinary noise takes it out before the idea has room to work. And critically, the strategy is negative *gross* — this is not a fee problem we can size our way out of.

By the level that was swept

LEVEL	TRADES	NET
DAY_HIGH	15	₹-218.10
ROUND	11	₹-144.80
PDL	3	₹-84.18
DAY_LOW	2	₹-54.24

14:15 – every agent was watching a price frozen at 9:16am

That "day high" row prompted a check. Each agent's levels are computed once, when it is *assigned* a stock. An agent that keeps its stock all session **never refreshes them**.

So "today's high" actually means *the high of the first minute of trading*, held for six hours.

STOCK	LEVEL AT 9:16	ACTUAL LEVEL LATER	DRIFT
INDSWFTLAB	337.80	372.00	10.12%
DEEPINDS	651.95	671.90	3.06%
GOKEX	800.00	805.25	0.66%
MINDACORP	723.50	725.00	0.21%

Why this probably explains the loss

A stale level is not a level anyone is watching. The entire premise of a sweep is that *other people's orders rest there* — and nobody has orders parked at a price that stopped being significant at 9:17am.

Thirteen of the twenty-five trades fired on exactly that. This is the single most likely cause of the negative gross, and it is a **bug rather than a strategy failure** — which means today's loss is **not yet a fair test of the idea**.

Not fixed today, deliberately. Changing the measuring instrument mid-measurement would have wasted the rest of the session. It goes in at tomorrow's open.

What needs your decision

Four open calls. Each is stated with options rather than a single recommendation, because these are yours to weigh.

1 - Real money – when?

My position is unchanged: not until the ₹2,000 top-up lands, the equity lane has produced cards on three separate days, and the floor completes one clean session. None of those is true yet.

- a Hold until all three gates pass — my recommendation
- b Go anyway at minimum size as a pure execution test, measuring fills against the card price. Defensible — but it needs a card to exist, and the equity lane has never produced one.

2 - Fix stale levels, then re-run

Refresh each agent's levels every few minutes instead of once at assignment. Today's loss was measured on a broken instrument; the number after this fix is the first honest test of the entry rule.

- a Fix at tomorrow's open and re-measure over a full clean session
- b Also widen the stop — 17 stops against 8 targets suggests it sits inside the noise band

3 - How far to trust the chart read

It found in two charts what the mechanical rule could not see at all. A strong first showing — but it is two charts.

- a Run it on every sweep signal for a week and compare its verdicts against what price actually did
- b Let it gate entries — nothing opens unless the chart read agrees. Slower, and it cannot run on every tick.

4 - The unresolved contradiction

Sweep-and-reclaim scored **+0.278pp** live but **-0.017%** in our historical falsification run. Both were measured properly. Either the live version picks up something the bar-level test flattened, or 118 examples is too small to trust. Worth resolving before anything else is built on it.