TradePilot

WEEKLY REVIEW · 3–7 AUGUST 2026

TP-WK-32 · SOUMYA SWAIN · CONFIDENTIAL

What this covers

Five trading sessions, every engine, every rupee — measured from the trade records rather than recalled. This is the first week TradePilot is treated as a product in its own right and the first with a revenue target attached, so the closing sections say plainly what that target requires and how far the system is from it.

Sessions	5 (Mon 3 – Fri 7 August)
Engines run	18 distinct, 11–17 on any given day
Trades placed	3,526
Net result	**−₹17,162**
Changes shipped	50
Findings that changed the plan	5

> **A note on the number.** The −₹17,162 is after trading costs, with costs corrected for the engines that have never booked any. Two experiments started mid-week and have only two sessions behind them; their figures are shown but should not be read as results.

At a glance

```
FOUND — AND IT CHANGES EVERYTHING
  The strategy's entry timing is measurably worse than picking a
  random moment in the same stock on the same day.

THE REAL SHAPE OF THE PROBLEM
  There IS an edge. It earns 0.069% per trade. Costs take 0.120%.
  We are not broken — we are smaller than the toll.

RUNNING
  Order-book collection began Friday. First real data in hand.
```

Where the week went

DAY	TRADES	NET	WHAT HAPPENED
Mon 3	539	**+₹17,273**	Best day of the week

DAY	TRADES	NET	WHAT HAPPENED
Tue 4	478	−₹21,291	Gave it all back
Wed 5	853	−₹16,903	Losses before costs, on a flat market
Thu 6	817	+₹5,628	
Fri 7	839	−₹1,869	

The primary engine lost ₹933 across the week. In isolation that reads badly; against the preceding eight weeks – which ran −1,894, +1,377, −7,575, −4,634, +805, −6,233, +681 – it is an ordinary week, not a bad one. That consistency is itself the finding: the system has not been getting worse. It has been flat-to-negative since June and this week did not change that.

What was found

The entry signal is worse than a coin flip

Holding the stock, the day and the stop-loss constant and changing only *when* the position was opened, a randomly chosen moment beat our timed entry by 0.18% per trade. Across five independent random draws, random won all five; giving those entries our own stop-loss did not close the gap. This is not a weak signal – it is one that fires reliably after the move it was meant to catch, a median of 126 minutes late, with only 0.76% of the day's range still available.

And the reason is structural

Every one of the six inputs the strategy scores on describes what has already happened: relative strength means the stock already outperformed, the breakout measure only fires once the range is broken, and the institutional-flow figure is published at the end of the day, so today's reading is yesterday's information. Nothing in the score looks forward. The 126-minute lag is not a defect in the code – it is the arithmetic of that input set.

The edge is real. It is just smaller than the toll

This is the week's most useful correction. The strategy earns **0.069% per trade** before costs, with a 47% win rate against a 42% break-even – genuinely above water. Trading costs take **0.120%**. The toll is 1.75 times the edge.

"No edge" and "edge smaller than the toll" are different problems with different answers, and the system had been diagnosed as the first. It is the second.

Holding longer does not fix it

The obvious response – pay the toll once, capture a bigger move – was tested against every entry since June. Held three days the trades looked profitable; subtracting what the market did over the same window, the profit was **zero to four decimal places**. Not skill, but being invested while the index rose 4.1%.

What was built

Fifty changes shipped. The substantial ones: the universe was audited against the exchange's lists and four dead symbols removed; the full 4,353-stock market was made browsable while the engines stay restricted to what they can fill; a portfolio view with a per-stock ledger went live; and the data feed moved to the licensed broker connection, 22 times faster and agreeing with the old source to the paisa. Three infrastructure faults were also fixed, each of which had been reporting success while doing nothing – the launch gate blocked the fleet from an open, the dashboard served an eighteen-day-old market, and the order-book collector ran two days producing zero bytes while recording success every time.

The ₹1 crore target

The target is ₹1 crore in a year from product revenue and trading profit together, with trading returns staged – 30% to begin, then 40%, then 50%, and capital scaled up once 70% is cleared. Measured against where the engines actually are:

RETURN GOAL	NEEDED PER DAY	NEEDED PER TRADE	VERSUS TODAY
30%	0.105%	₹23.52	4.6× the current gross edge
40%	0.135%	₹30.17	5.4×
50%	0.162%	₹36.36	6.2×
70%	0.212%	₹47.60	7.5×

Today the primary engine returns **−0.031% per day**, or −7.7% annualised. The first milestone is therefore not 30% – it is **zero**. Every multiple above assumes the gross edge first clears the cost, and it does not yet.

On the revenue side the arithmetic is friendlier and does not depend on the strategy working: ₹1 crore is roughly 835 subscribers at ₹999 a month held for a year, or 278 at ₹2,999. That is a distribution problem rather than a trading one – with the condition that the product must describe honestly what it does, which today means not selling a signal that loses to a coin flip.

What to be aware of

The strategy has produced no measurable skill at any horizon tested – intraday or multi-day, long or short. Four findings reported during the week collapsed when re-tested, every one because the measurement used information the system would not have had at the time; that pattern is now a standing check rather than a lesson. The live-money plan is on hold – it existed to promote whichever engine won a head-to-head, and every engine has lost money since June. One dependency remains manual: the broker connection expires at 6am daily and must be renewed by hand, and a missed renewal now costs a day of order-book collection.

Next week

1. **Protect the collection.** Friday produced the first real order-book data, 65 MB. Two to three weeks are needed before it can be tested, and it silently collected nothing for two days before that was caught. Verify daily.
2. **Test the two cheap ideas left** – whether index futures move before their constituents, and whether an overnight gap predicts the day. Both use data held.
3. **Leave the engines alone** – they are the out-of-sample record – and **start the product track**, which does not depend on the edge working.

Evidence · 3,526 trades, 18 engines, 5 sessions, costs 12 bps and corrected where unbooked · 50 commits · entry-vs-random 5 seeds t 2.76–4.24 · timeframe 1,924 entries vs same-window index · liquidity validated on 18,053 closed trades · from trade records, 9 Aug. Soumya Swain · soumya@suryaai.co.in