

ALGORITHM DIAGNOSIS REPORT

TradePilot v3 to v4 Algorithm Rebuild

Why the Current Model Fails — And How v4 Fixes It

Date: April 8, 2026

Project: TradePilot — Intraday Trading Signal Engine

Author: Soumya Swain

Version: v1.0

Table of Contents

1. Executive Summary

2. The Problem — v3 Algorithm Diagnosis

3. Root Cause Analysis

4. Today's Evidence: April 8, 2026

5. v4 Architecture — The Fix

6. Composite Scoring System

7. Implementation Roadmap

8. Success Criteria & Validation Plan

1. Executive Summary

The Bottom Line: TradePilot's v3 model is fundamentally misaligned with its purpose. It was built to predict 5-day forward returns (a mean-reversion strategy) but is being used for **intraday momentum trading**. The result: 96% of stocks are flagged AVOID, only 2 BUY signals per day, and 80%+ of profitable intraday moves are missed.

On April 8, 2026, 15 out of 50 Nifty stocks gained more than 1% intraday. The model identified only **2 of them**. The remaining 13 stocks — representing 24.1% combined upside — were marked AVOID. This is not a calibration problem. It is an **architecture problem**.

The v4 Solution: A complete rebuild around 7 signal layers (ML + Relative Strength + ORB + VWAP + FII/DII + Options OI + Volume), percentile-based ranking (always 10 BUY picks), and real NSE institutional data replacing simulated options chains. Estimated effort: 39 hours across 7 phases (~2 weeks).

2. The Problem – v3 Algorithm Diagnosis

2.1 Score Distribution is Catastrophically Skewed

The v3 scoring engine produces a distribution with **mean = 17** and **median = 12.6** on a 0–100 scale. With BUY thresholds set at 55+ (v2) and 50–60 (v3 regime-dependent), the vast majority of stocks can never reach the BUY zone regardless of market conditions.




Figure 1: v3 composite score distribution — 96% of stocks trapped below the BUY threshold

2.2 The Fundamental Mismatch

Aspect	What v3 Does	What We Need
Prediction Target	5-day forward return > 0.5% (binary)	Intraday return magnitude (regression)
Strategy Type	Mean-reversion (buy dips)	Momentum (buy strength)

Time Horizon	Hold 5 days	Entry 9:35 AM, exit by 3:15 PM
Signal Method	Absolute score threshold	Relative ranking (always top 20%)
Data Sources	Daily OHLCV only	Intraday candles + FII + Options OI
Options Data	SIMULATED (np.random)	Real NSE options chain

Critical Finding: The options chain data in v3 is entirely simulated using `np.random`. Every options-based signal (PCR, max pain, OI buildup) in the current system is noise, not signal. This alone invalidates any options-based scoring component.

3. Root Cause Analysis

We identified 6 root causes contributing to the model's failure to capture intraday opportunities. The waterfall chart below shows each cause's contribution to missed signals:




Figure 2: Root cause impact waterfall — wrong prediction target is the #1 driver (35%)

3.1 Wrong Prediction Target (35% impact)

The model's training target is: *"Will this stock return > 0.5% over 5 days?"* (binary classification at `trading_engine_v3.py:213-226`). For intraday trading, we need: *"How much will this stock move between 9:30 AM and 3:15 PM today?"* (regression). This is not fixable by tuning — it requires retraining with a fundamentally different target variable.

3.2 No Intraday Features (25% impact)

v3's top features are all daily/weekly indicators designed for swing trading:

v3 Feature	Weight	Type	Useful for Intraday?
------------	--------	------	----------------------

ATR (14-period)	13.25%	Volatility range	No
MACD histogram	7.56%	Trend momentum	Marginal
Volatility 20d	6.69%	Historical vol	No
Pct from 52W low	6.21%	Mean-reversion	No
RS 20-day	5.78%	Relative strength	Partially

Missing entirely: VWAP position, Opening Range Breakout (ORB), gap analysis, first-hour momentum, volume surge detection.

3.3 Simulated Options Data (15% impact)

The options chain data fed into the model is generated via `np.random`, making every options-derived signal (PCR, max pain, OI buildup) pure noise.

3.4 Absolute Thresholds (12% impact)

v3 uses fixed thresholds (score ≥ 55 for BUY). In a low-volatility market, no stock crosses 55. In a bull run, dozens might. The threshold doesn't adapt, causing:

- Bull days: too many BUY signals (no relative ranking)
- Normal/bear days: zero to two BUY signals (today's problem)

3.5 Poor ML Precision (8% impact)

Training precision of 44.79% means the model's positive predictions are wrong more often than right. With 56.46% recall, it catches some winners but generates too many false positives in the rare cases it does signal BUY.

3.6 No FII/DII Flow Data (5% impact)

Foreign Institutional Investor (FII) and Domestic Institutional Investor (DII) net flows are the single best leading indicator of broad market direction on

NSE. v3 ignores them entirely.

4. Today's Evidence: April 8, 2026

The chart below shows every Nifty 50 stock that gained >1% today, along with the v3 model's signal and score for each:




Figure 3: April 8 intraday winners — model correctly identified only 2 out of 10 (20% capture rate)

Key Observation: SHRIRAMFIN (+3.20%) received a score of just 6.5, while ONGC (+1.44%) got 58.7. The model's score bears little correlation with actual intraday performance. TITAN was the only stock where both the model and reality aligned strongly (score 67.3, actual +3.97%).

By the Numbers

Metric	Value
--------	-------

Stocks up > 1% today	15
Correctly flagged BUY	2
Flagged HOLD (missed)	1
Flagged AVOID (missed)	12
Combined missed upside	24.1%
Capital deployed	Rs 2L of 10L (20%)
Capital sitting idle	Rs 8L (80%)

5. v4 Architecture – The Fix

5.1 Feature Coverage: v3 vs v4

The radar chart below shows how v4 dramatically expands signal coverage across 8 critical dimensions:




Figure 4: Feature coverage radar — v3 (red) covers only ML + basic volume; v4 (green) covers all 8 dimensions

5.2 The Three Design Pillars of v4

Pillar 1 — Regression, Not Classification: v4 predicts the *magnitude* of intraday return using LightGBM regression. Instead of "will it go up? yes/no", we predict "by how much?". This preserves ranking information that binary classification destroys.

Pillar 2 — Relative Ranking, Not Absolute Thresholds: Every day, all 50 stocks are ranked by composite score. The top 20% (10 stocks) = BUY, next 30% = HOLD, bottom 50% = AVOID. This **guarantees 10 BUY signals daily** regardless of market conditions.

Pillar 3 — Real Institutional Data: FII/DII flows, real options chains (via nsepython), VWAP from NSE quotes, and intraday candles replace all simulated data. Every signal layer uses verified market data.

5.3 v4 Feature Architecture (19 Features)

Layer	Features	Count	Source
Daily Context	Nifty change, FII/DII flows, sector, VIX, volume ratio, breadth, prev range	9	nsepython, yfinance
Intraday	VWAP deviation, ORB signal, ORB range, intraday trend, volume surge	5	nsetools, yfinance
Institutional	Put-Call Ratio, max pain distance, OI buildup signal	3	nsepython (real OI)
Relative Strength	5-day RS vs Nifty, 20-day RS vs Nifty	2	Computed from prices

Total	19
-------	----

6. Composite Scoring System

6.1 Weight Distribution




Figure 5: v4 composite score is a weighted blend of 7 signal layers — no single layer dominates

6.2 How Each Signal Layer Works

Signal Layer	Weight	What It Captures	Why It Matters
--------------	--------	------------------	----------------

ML Prediction	25%	Predicted intraday return magnitude	Pattern recognition across 19 features
Relative Strength	20%	5d + 20d price performance vs Nifty	Momentum stocks outperform — well-documented factor
ORB Breakout	15%	First 15-min range breakout direction	60-89% win rate on NSE (proven strategy)
VWAP Position	10%	Price vs Volume-Weighted Average Price	Above VWAP = institutional buying pressure
FII/DII Flow	10%	5-day rolling net institutional flow	Smart money direction — leading indicator
Options OI	10%	PCR + OI buildup pattern	Options writers' positioning reveals conviction
Volume	10%	Today's volume vs 20-day average	Confirms moves — breakouts on low volume fail

6.3 Signal Generation: v3 vs v4




Figure 6: v3 produces 2 BUY / 46 AVOID; v4 guarantees 10 BUY / 15 HOLD / 25 AVOID daily

The Key Design Fix: Relative ranking (always top 20% = BUY = 10 stocks) instead of absolute thresholds (score ≥ 55 = only 2 stocks). This means v4 *always* picks the best available opportunities, even in a flat market.

7. Implementation Roadmap

7.1 Timeline Overview




Figure 7: 7 phases, 39 hours total — Phases 1+2 run in parallel (Days 1-3)

7.2 Phase Details

Phase	What	Hours	Key Output
1	Real NSE Data Pipeline	8	FII/DII, options OI, VWAP, intraday candles
2	Intraday Feature Engineering	7	ORB, VWAP position, gap, momentum, volume surge
3	ML Regression Engine	7	LightGBM predicting return magnitude (not binary)
4	Composite Scorer + Ranking	5	7-layer weighted score + percentile classification
5	Kelly Position Sizing + Paper Trade	5	Confidence-based allocation, v3 vs v4 parallel run

6	Dashboard Integration	4	?engine=v4 toggle, daily comparison reports
7	Real-Time 10-min Scanning	3	Live signal updates during market hours
Total		39	~2 weeks at 3-4 hrs/day

7.3 Dependency Graph

Parallel Opportunity: Phases 1 (data pipeline) and 2 (feature engineering) have no dependencies on each other and can run simultaneously, saving 3 days. Phases 5 and 6 can also partially overlap after Phase 4 completes.

7.4 Data Flow (Market Day)

Time	Action	Module
9:00 AM	Pre-market: fetch FII/DII, options chain	data_nse.py
9:15 AM	Market open: gap analysis, first quotes	data_nse.py
9:30 AM	ORB forms (15-min candle completes)	features_intraday.py
9:35 AM	First scoring run → deploy capital into top picks	composite_scorer.py
9:45+	Re-score every 10 min, react to signal changes	composite_scorer.py

3:15
PM

Force close all positions

v4-paper-trade.py

3:30
PM

Compare v3 vs v4 vs actual → daily report

v4-daily-
compare.py

8. Success Criteria & Validation Plan

8.1 Performance Targets




Figure 8: Performance gap between v3 current state and v4 targets across 4 key metrics

Metric	v3 (Current)	v4 (Target)	Improvement
BUY signals per day	2	8–12	5x increase
Hit rate (BUY stocks end green)	~50%	>55%	+5 percentage points
Avg return of top-10 picks	Unknown	>0.5% intraday	Measurable for first time
Missed opportunities (>1%, not BUY)	13 of 15	<5 of 15	60% fewer misses
Capital deployed	Rs 2L of 10L (20%)	Rs 8L+ of 10L (80%)	4x more capital working

8.2 Validation Protocol

- 1. Parallel Run (2 weeks):** v3 and v4 score all 50 stocks simultaneously.
Paper trade Rs 5L with each engine.
- 2. Daily Comparison:** At 3:30 PM, compare v3 picks vs v4 picks vs actual returns. Log to `docs/validation/v4-comparison/`.
- 3. Phase Gates:**
 - After Phase 1: `get_fii_dii_daily()` returns real data (not mock)
 - After Phase 2: ORB computation shows breakout signal on today's data
 - After Phase 4: `score_all_stocks()` returns 10 BUY signals (not 2)
 - After Phase 6: Dashboard serves v4 scores at `/api/scores?engine=v4`
- 4. Go/No-Go:** After 2 weeks, if v4 hit rate > v3 hit rate AND missed opportunities < 5, swap v4 as primary engine.

8.3 Risk Mitigation

Risk	Likelihood	Mitigation
NSE API rate limits / blocks	Medium	Local caching per day, staggered requests
Options data unavailable for non-F&O stocks	High	OI features only for F&O subset; fallback to volume-only for others
LightGBM overfits on 3-month data	Medium	Walk-forward validation, out-of-sample MAE tracking
v4 underperforms v3	Low	Keep v3 running; switch back if v4 fails 2-week test

Conclusion: The v3 model was built for the wrong problem (mean-reversion over 5 days) and uses the wrong data (simulated options, no intraday signals). v4 is designed from the ground up for intraday momentum trading with real institutional data, 7 signal layers, and a ranking system that guarantees actionable picks every market day. The parallel build approach means zero disruption to the live dashboard during development.