

- [TradePilot – Master Roadmap](#)
 - [0. The one-paragraph state of play](#)
 - [1. Sequencing – do these in order](#)
 - [2. Track A – Make it profitable \(from root-cause work\)](#)
 - [3. Track B – Autonomy upgrade \(T3 → T4\)](#)
 - [4. Track C – Competitive moat / product](#)
 - [5. Track D – Compliance \(the SEBI answer\)](#)
 - [6. Track E – Validation & risk](#)
 - [7. Track F – Research follow-ups](#)
 - [8. The honest bottom line](#)
 - [9. Track G – AGENTIC TRADING PLATFORM \(productization\)](#)
 - [10. Progress log – 2026-06-30](#)

TradePilot – Master Roadmap


Everything we now know we need to do – performance, autonomy, moat, compliance, validation, research

Project	TradePilot (autonomous agentic trading)
Document	Master roadmap (consolidates root-cause + competitive + compliance research)
Version	v1.0.0
Status	Active
Created	2026-06-28
Sources	Root-cause investigation (11 agents) · competitive deep-research (105 agents) · SEBI + Tier-2 + Tier-4 research (3 agents)
Author	Soumya Swain
Email	soumya@sidewall.in

0. The one-paragraph state of play

TradePilot is a **T3 (Conditional Autonomy)** agentic trading system – the same autonomy level credible institutional operators actually run at, and ahead of the unverified “fully autonomous” marketing of pre-launch startups. It has **no direct competitor in India** (every Indian platform is a rule-based strategy builder, not agentic) and the research literature **validates its core design choice** (execution discipline over LLM-alpha, which is largely lookahead-bias illusion). The gaps are the same ones almost every peer shares: **live capital, an audited track record, and scale**. Nothing it does today is non-compliant – paper trading is outside SEBI’s framework entirely. The roadmap below sequences the work so we **prove profitability first, go live cheaply second, and only take on heavy compliance if we productize**.

1. Sequencing – do these in order



roadmap_2026-06-28-diagram-0

The rule: don't spend a rupee or a day on heavy (product) compliance until the agent is *proven profitable* and *you've decided to sell it*. Paper-perfect first.

2. Track A – Make it profitable (from root-cause work)

ID	ACTION	STATUS
RC-1	Long-only shadow (<code>v5_long</code> , NIFTY-200) racing live v5	Live Mon 06-29
RC-2	If long-only wins ~2 wks → gate shorts in live v5	Pending RC-1 result
RC-3	Fix entry timing (collapse SHORT_BLOCK lag, 2nd deploy pass)	Queued
RC-4	Concentrate capital (cut cap 20→5-8; deploy more of ₹10L)	Queued
–	Data-stall fix (yfinance per-process cache)	Done
–	Engine roster trimmed 6→4 (lean stack)	Done

Gate: RC-1's two-week verdict (auto-Telegram'd daily at 15:40) decides whether the whole "short book is the bleed" thesis holds. Everything downstream waits on it.

3. Track B – Autonomy upgrade (T3 → T4)

We sit at T3: AI runs the full intraday loop autonomously, humans approve strategy changes. The peers ahead of us (Self-Driving Portfolio's meta-agent) reach higher by **self-modifying strategy parameters**.

ID	ACTION	PRIORITY
AU-1	Let the learnings-store <i>propose</i> strategy-param changes (regime weights, signal set, risk limits) as PRs for human approval	Medium
AU-2	Multi-agent debate layer (bull/bear, à la TradingAgents) before signal commit	Low / optional
AU-3	Eventually: auto-apply low-risk param changes without sign-off → T4	Future (only after live track record)

Note: T4 (self-modifying without approval) is *not* a near-term goal – it's only sensible once the system has a proven live track record. T3 is the right level for now and matches what real funds run.

4. Track C – Competitive moat / product

AREA	OUR POSITION	ACTION
Autonomy	Uncontested in India (category gap)	Lead with it – “agent that decides what to trade,” not another strategy builder
Live execution	Incumbents far ahead (Tradetron 1.5M trades/mo, 60+ brokers)	Build a real execution layer (Kite order API) – the biggest product gap
Users / distribution	None vs Sensibull 1M+, AlgoTest ~25k	Defer until agent is proven; distribution is a post-PMF problem
Positioning	Different <i>category</i> , not a better builder	Target the user who doesn't know <i>what</i> to trade, not the one who needs execution infra

5. Track D – Compliance (the SEBI answer)

Paper trading = zero obligations. SEBI's framework (Feb 2025 circular, fully in force **April 1, 2026**) regulates *orders on exchanges via broker APIs*. We place none. Nothing to fix today.

PHASE	TRIGGER	OBLIGATIONS	BURDEN
D-0 (now)	Paper only	None. (Document the no-order-routing architecture for the record.)	Zero
D-1	Go live, own capital via Kite	Static IP + Zerodha whitelist · software kill-switch · 5-yr order audit log · confirm Kite algo TOS · OPS profiling (we run ~0.0024 OPS avg, peak <1 – far under the 10-OPS registration line) · family-account carve-out only	Light – a checklist, not a minefield
D-2	Offer to others (product)	SEBI Research Analyst license (we're black-box) · NSE/BSE empanelment · broker partnership · ISO 27001:2022 · biannual CERT-In VAPT · per-strategy algo registration · grievance redressal	Heavy – 6-12 months + real spend

Comply-before-live checklist (Phase D-1): static IP → Zerodha whitelist → kill-switch in code → order audit log (5-yr) → Kite Connect algo-TOS review → OPS peak instrumentation → written algo design doc. None are blockers; all are a weekend of work plus broker setup.

Hard line: running the algo for anyone outside your family (self/spouse/dependent parents/children) makes you an “algo provider” → instantly triggers the heavy D-2 tier. Keep it personal until you deliberately productize.

6. Track E – Validation & risk

ID	ACTION	WHY
V-1	Keep alpha in execution, not LLM signals	Literature: LLM alpha is largely lookahead-bias; collapses out-of-sample
V-2	Define the paper→live trust gate (e.g. N profitable weeks, max-drawdown bound)	Don't risk real capital until earned
V-3	Build toward an audited track record	The universal gap vs all peers; the thing that makes claims credible
V-4	Stress-test robustness (perturbation)	TradeTrap: agentic systems are fragile; ours must not be

7. Track F – Research follow-ups

ID	QUESTION	PRIORITY
R-1	Numerai deep dive – closest verified AI-native fund (JPMorgan committed up to \$500M, ~25% net 2024); how “AI-driven” is it really?	Medium
R-2	SEBI RA-license process detail (net worth, timeline, cost) – the gate for D-2	Medium (if productizing)
R-3	Jarvis Invest reverse-engineer – closest Indian regulatory analog (SEBI RA + AI portfolio, ₹500cr+)	Medium
R-4	Zerodha Kite Connect algo-order TOS – does own-account automation below 10 OPS need separate approval?	High (before D-1)
R-5	Standard Signal SEC EDGAR / Form ADV check – is it actually a registered live fund?	Low (curiosity)

8. The honest bottom line

- **We are not behind.** At parity on architecture, ahead on operational maturity, uncontested on India, at the same autonomy level (T3) as credible real funds.
- **We are unproven** on the three things that matter for credibility: live P&L, audit, scale – but so is almost everyone else.
- **The path is clear and de-risked:** prove profitability (Track A) → go live cheaply for own capital (D-1, light) → only then consider product + heavy compliance (D-2/C).
- **Do not** chase LLM-alpha sophistication or heavy compliance prematurely; the moat is autonomy + discipline + (eventually) a real track record.

9. Track G – AGENTIC TRADING PLATFORM (productization)

The destination: turn the single-user autonomous engine into a genuine agentic platform. Product shape (research-recommended for India): **“BYOB White-Box Agent”** – your own autonomous AI agent running in your own Zerodha/Fyers account; set a mandate once; it researches, decides, trades; you see every reasoning step. Platform holds **zero capital** (no PMS licence); **white-box** sidesteps the heavy black-box SEBI RA gate; the explainable decision log is simultaneously the trust UX + the compliance audit trail.

PHASE	WHEN	BUILD	COMPLIANCE
G-0 • Personal T4	now → 6mo	Decompose monolith → 9-role agent DAG (Planner · Orchestrator · Alpha · Risk-gate · Portfolio · Backtest · Execution · Audit · Memory); meta-agent rewrites prompts; UUID memory with train/eval namespace separation (structural anti-overfit); explainability log	Zero (paper/own capital)
G-1 • Invite-only BYOB	6-12mo	3-10 trusted users; per-user IPS (mandate) + broker-credential vault + static IP; white-box algo-ID; per-user audit namespace	Light (white-box, no RA licence)
G-2 • Closed beta	12-18mo	50-200 users; subscription; copy-agent of our own variants; full audit infra	SEBI RA licence, ISO 27001 prep
G-3 • Platform / marketplace	18-24mo	Third-party agent SDK; per-agent algo-ID + track record	Empanelment, VAPT, ISO 27001

Key architectural unlock: the agentic DAG (alpha agents blind to evaluation-window returns) is the *same* investment as the overfitting fix in Track A – platform and profitability converge. **Gate before any of this:** the conviction→P&L validation (instrumented 2026-06-30), and Track A profitability proven. Reference: [STRATEGY_RESEARCH_2026-06-30](#) Pillar C.

10. Progress log – 2026-06-30

ITEM	STATUS
Conviction logging on closed trades (all engines)	✅ Built – unblocks dynamic-allocation validation in ~5 sessions
v5_flip fast-regime-flip shadow + wired into rotation (roster=5)	✅ Built – in A/B testing from next launch
Fast-regime-flip design (validated 4x vs data)	✅ Spec: DESIGN_fast-regime-flip_2026-06-30
Red-day watchdog (live loss attribution)	✅ Running
Dynamic per-stock-trend allocation + net-exposure cap	🕒 Designed; gated on conviction data + Track A
Agentic platform (Track G above)	🕒 Roadmapped; gated on profitability proof

Next decision point: EOD review of [v5_flip](#) vs the roster on the next red day – does fast activation of the 8/12 tilt cut red-day losses without false-triggering on green days?