

- [TradePilot – Session Changelog \(June 2026 arc\)](#)
 - [A. Root-cause investigation – why we were losing](#)
 - [B. Debugs / data-integrity fixes](#)
 - [C. Engine roster consolidation](#)
 - [D. Competitive & strategy research](#)
 - [E. Red-day / regime-flip work \(06-30\) – Sarathi-disciplined validations](#)
 - [F. Features / tools built](#)
 - [G. Learnings persisted \(DevPilot store\)](#)
 - [H. Project memory](#)

TradePilot – Session Changelog (June 2026 arc)

Complete record: every finding, debug, fix, feature, validation, research output, and design – with the doc + learning where each lives.

Project	TradePilot (autonomous agentic trading)
Document	Master session changelog / index
Version	v1.0.0
Period	2026-06-22 → 2026-06-30
Status	Living index – append new work here
Author	Soumya Swain
Email	soumya@sidewall.in

A. Root-cause investigation – why we were losing

FINDING	VERDICT
The edge is execution discipline, not ML/AI	On the 82%-WR original, winner score $\approx$ loser score; ML IC 0.006
Complexity diluted the edge	Win-rate 82% $\rightarrow$ $\sim$ 48% as trades/day 17 $\rightarrow$ 45 (NIFTY-50 $\rightarrow$ 200, +shorts, 4 pools)
<code>FLAT_FORCE_EXIT</code> – exonerated	Trade-level audit: it <i>helped</i> (+₹30, +₹299) on the days it was blamed
Late entry is the day-to-day bleed	v5 enters winning names 30–60 min late (rescore loop) $\rightarrow$ booked as STOPLOSS not TARGET
Over-trading = cost drag	Gross +₹4,691 but cost –₹7,153 = net –₹2,462 over 8 days; excess churn = 205% of net
The SHORT book is the net bleed	Longs +₹1,149 vs shorts –₹3,611 (06-16..25); shorts hit TARGET only 4.1% (longs 8.5%)
v5_classic (frozen original) beats live v5	+₹111,803 vs +₹116,357 over 45 sess, but v5's lead is a 2-outlier-day mirage; June: classic +₹12,891 vs v5 +₹690 (18 $\times$)
OLD 5-tree model (sibling A/B)	+₹122k headline is a leverage(313%)+gross+72%-from-4-days mirage – BUT the 5-tree beating the 1,735-tree is real overfitting evidence

Docs: `ROOT_CAUSE_ANALYSIS_2026-06-24.{md,pdf}` , `ROOT_CAUSE_ANALYSIS_II_2026-06-26.{md,pdf}` .

B. Debugs / data-integrity fixes

BUG	ROOT CAUSE	FIX
06-26 all-day data stall (engines alive, 0 trades)	yfinance shared SQLite tz-cache contention under concurrent multi-process load $\rightarrow$ <code>OperationalError: unable to open database file</code> $\rightarrow$ downloads return None	Per-process cache isolation (<code>yf.set_tz_cache_location</code> keyed by <code>ENGINE_NAME/PID</code>) in <code>data_nse.get_yfinance()</code> + <code>missed-opportunities-watchdog.py</code> ; cleared corrupt shared db.  verified live download works
Conviction not stored on closed trades	<code>close_position</code> rebuilt the record from a hardcoded subset, dropping <code>score</code> / <code>reasons</code> that were on the open position	Carry score/direction/reasons/sl_price/target_price/entry_date/trailing_activated into closed record across all engines (v5 family + v5_classic + v7_regime)

C. Engine roster consolidation

- **6 $\rightarrow$ 4 lean roster** (06-26 audit): retired `v5_nom1` (redundant – ran v5 twice), `v5_apr` (tracked v5 +₹78/9d), `v7_regime` (flat + WFO-negative). State preserved (commented, not deleted).
- **+ v5_long** (RC-1 long-only) and **+ v5_flip** (fast regime-flip) $\rightarrow$ **roster now 5**: v5 · v5_classic · v5_long · v5_cut · v5_flip. Consistent across `launch-market.sh` , `crash-watchdog.sh` , `engine-compare.py` .

D. Competitive & strategy research

OUTPUT	VERDICT
Competitive landscape (105-agent deep-research)	No direct competitor. Peers are ~95% research prototypes; only commercial one (Standard Signal) unverified
Autonomy benchmark	TradePilot = T3 (Conditional Autonomy) – same level credible real funds run at
Tier-2 India	Zero Indian platforms are agentic (all rule-based builders) – category gap
Tier-4 funds	Every real-money fund keeps a human in the strategy loop; Numerai closest verified AI-native
SEBI compliance	Paper = zero obligations; live own-capital = light ; product = heavy (RA licence, empanelment, ISO 27001)
3-pillar strategy	Profitability + UX + platform all share one spine: a simple, explainable, clean-execution decision pipeline

Docs: [COMPETITIVE_LANDSCAPE_2026-06-28.{md,pdf}](#) , [STRATEGY_RESEARCH_2026-06-30.{md,pdf}](#) , [ROADMAP_2026-06-28.{md,pdf}](#) .

E. Red-day / regime-flip work (06-30) – Sarathi-disciplined validations

Live red day (NIFTY –0.73%) → longs bled, shorts green → built diagnosis + design. **The data corrected coarse framings four times:**

CLAIM TESTED	RESULT
"No flips after 13:30"	REFUTED – v5 post-1pm entries +₹30,515 vs morning –₹2,724. Rule removed; flip is bidirectional + all-session
"short=red, long=green" (flip the book)	REFUTED – edge is stock-selection; longs profit on DOWN days (+ ₹79,428, 56% green); shorts profit on their own; COALINDIA shorted ×6 = +₹23,197. → tilt the ratio, don't flip
Tilt magnitude / trigger	Existing BEAR 8/12 is right ; trigger is hard-down < –0.6% (mild-down still favours longs), NOT –0.5%; don't chase all-in (linear-model + reversal + small-sample risk)
Fixed ratio vs dynamic	Engine does not adapt (short-share flat ~45% across up/down days). Principled design = per-stock trend direction + net-exposure risk cap , not a fixed ratio – but needs conviction data (now logging)

Docs: [DESIGN_fast-regime-flip_2026-06-30.md](#) .

F. Features / tools built

FEATURE	WHAT
<code>scripts/red-day-watchdog.py</code>	Real-time loss attribution (long vs short) + red-day counterfactual + Telegram alert
<code>scripts/engine-compare.py</code> + <code>launchd</code>	Daily 15:40 Telegram four/five-way engine scorecard (auto)
<code>docs/decision-dashboard.html</code> + <code>/decisions</code> route + pageswitch nav	Decision dashboard integrated into the live Flask dashboard (every page)
yfinance per-process cache fix	Data-stall root fix in <code>data_nse</code> + watchdog
Conviction logging	score/reasons/etc. on closed trades, all engines
<code>scripts/v5_flip-paper-trade.py</code> + wiring	Fast intraday regime-flip shadow (5-min tape → BEAR 8/12 tilt on confirmed hard-down, bidirectional)

G. Learnings persisted (DevPilot store)

~20 learnings stored this arc (search `dp lrc recall "tradepilot"`): root-cause, edge-source, over-trading, short-book, v5_classic, OLD-5tree, SEBI, competitive, autonomy-T3, 3-pillar, red-day, fast-flip design, 2nd-half validation, direction-vs-regime, tilt-magnitude, dynamic-allocation, conviction-logging + v5_flip build.

H. Project memory

`tradepilot-root-cause` , `tradepilot-competitive-roadmap` (+ MEMORY.md index).